

IDIAP 2015-03-17

OpenCL

Crash Course

frederic.dubouchet@idiap.ch

Sources + slides

http://github.com/anirul/OpenCL_Crash_Course.git

- Known to work on Linux/OSX:
 - a C++ compiler (g++/clang++)
 - the OpenCL Header + libs
 - cmake / OpenCV / Boost / GL / GLU / glut

Plan

- General overview (GPGPU -> OpenCL)
- Code dive (various examples)
- Conclusion (Optimisation tips)

General Overview

- GPGPU - Technology overview
- OpenCL
 - Language and API
 - Device Model
 - Objects

GPGPU

General Purpose GPU

- Using Graphical Processing Unit to compute.
 - Shader languages (GLSL / DirectX)
 - CUDA (Nvidia proprietary)
 - DirectCompute (Windows)
 - OpenACC (no free compiler support yet)

OpenCL

heterogeneous computing platforms

- Khronos (OpenGL, Vulkan, COLLADA, etc...)
 - Intel, QUALCOMM, AMD, Altera Corporation, Vivante Corporation, Xilinx, Inc., MediaTek Inc, ARM Limited, Imagination Technologies, Apple, Inc., STMicroelectronics International NV, ARM, IBM Corporation, Creative Labs, NVIDIA, Samsung Electronics.
- Work on CPU / GPU / DSP / FPGA ...
- Open Standard

OpenCL

Language and API



The diagram consists of two blue ovals stacked vertically. The top oval contains the text 'CPU' and the bottom oval contains the text 'GPU'. To the right of these ovals is a bulleted list of OpenCL features.

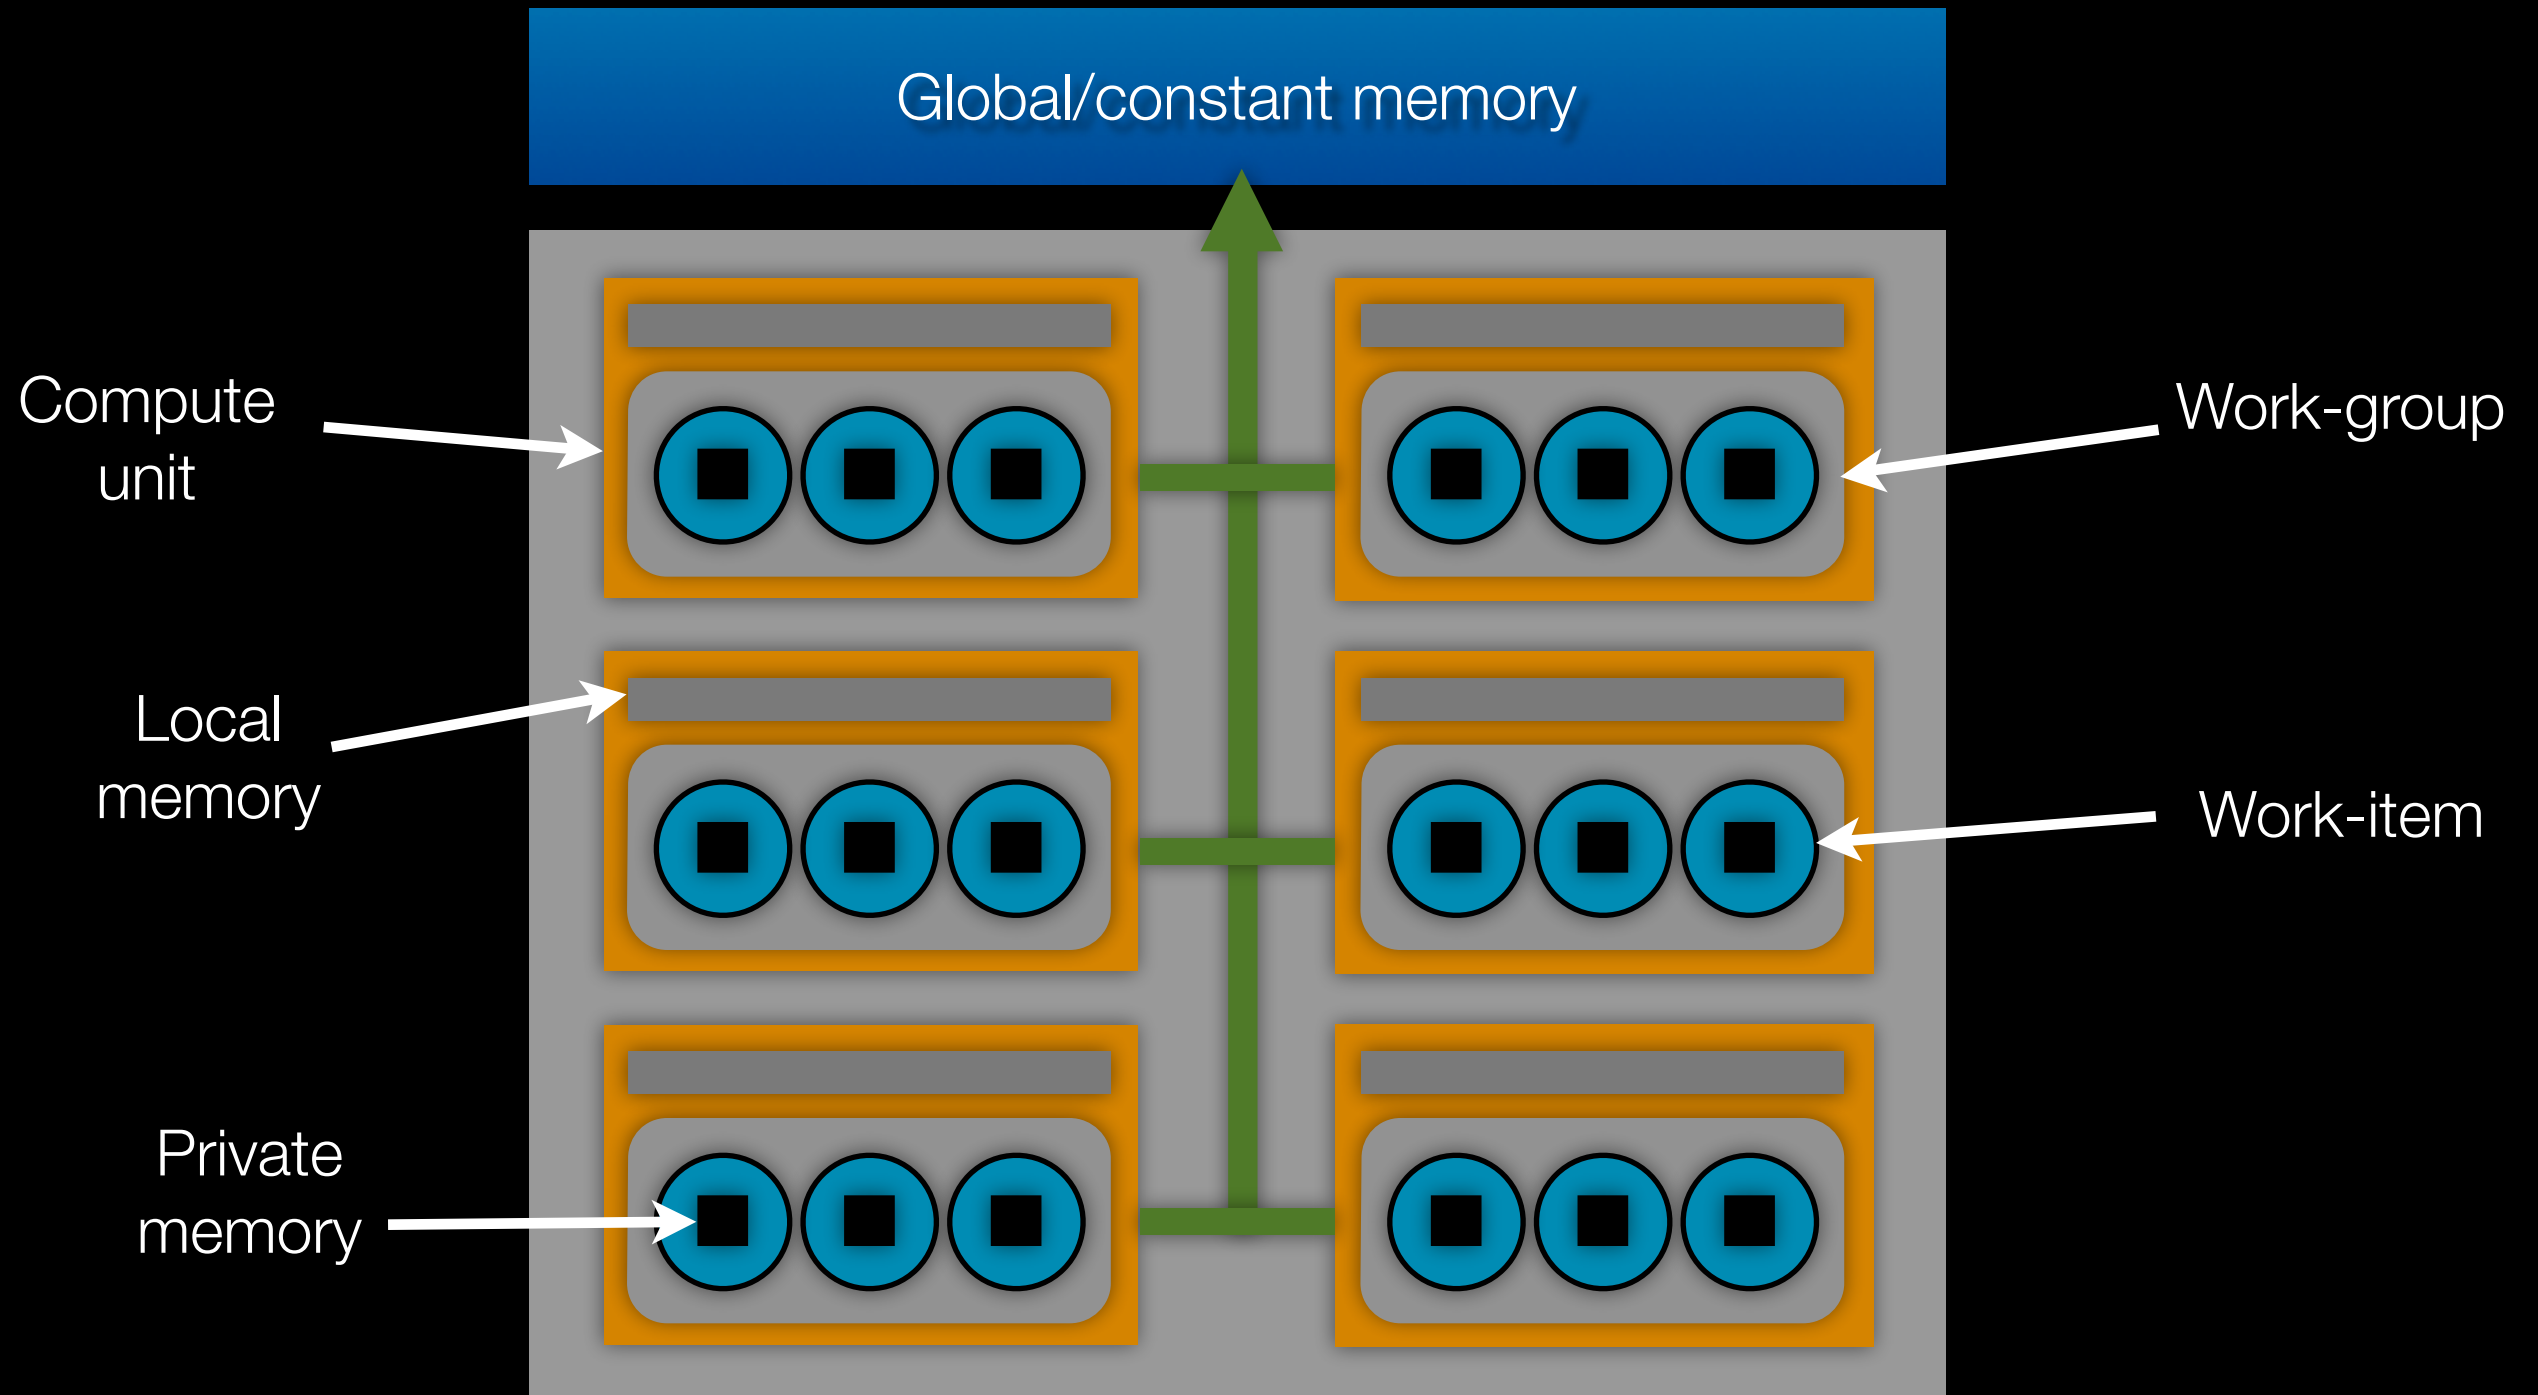
CPU

GPU

- ✧ An API (run on the Host)
 - ✧ in C but with interface to many other languages
- ✧ A language (run on the Device)
 - ✧ vector oriented
 - ✧ C99 inspired (2.1 -> C++14)


OpenCL

Device Architecture



OpenCL

Objects




Code Dive

- Khronos C++ wrapper
- Simple - (as it can be!)
- Floyd-Warshall - (memory coalescing)
- Histogram (reduce, local memory)
- Video (the full stack)

Khronos C++ wrapper

- ✧ Officially published by Khronos
- ✧ Template based
- ✧ Header only
- ✧ 100% portable
- ✧ Object Oriented



Simple

- Very simple example using the C++ wrapper
- Compute the product of 2 vectors
- Single file only OpenCL dependent
 - `simple.cpp`
- Not a practical example!

Device code

```
// very simple kernel

kernel void simple(
    global read_only float* in1,
    global read_only float* in2,
    global write_only float* out)
{
    const uint pos = get_global_id(0);
    out[pos] = in1[pos] * in2[pos];
}
```

- Simple product of two vectors

Host code (1)

```
#define __CL_ENABLE_EXCEPTIONS  
#include <CL/cl.hpp>
```

- Add exception support to OpenCL

```
std::vector<cl::Platform> platforms;  
cl::Platform::get(&platforms);  
std::vector<cl::Device> devices_  
platforms[platform_id].getDevices(  
    CL_DEVICE_TYPE_ALL,  
    &devices_);
```

- Get Platform and device

Host code (2)

```
cl_context_properties properties[] = {
    CL_CONTEXT_PLATFORM,
    (cl_context_properties) (platforms[platform_id])(),
    0
};
cl::Context context_ = cl::Context(CL_DEVICE_TYPE_ALL, properties);
cl::CommandQueue queue_(context_, devices_[device_id], 0, nullptr);
```

- Generate a context

```
cl::Program::Sources source(
    1,
    std::make_pair(
        kernel_source.c_str(),
        kernel_source.size())));
cl::Program program_(context_, source);
program_.build(devices_);
cl::Kernel kernel_(program_, "simple");
```

- Get the source build and select a kernel

Host code (3)

```
cl::Buffer buf_in1_ = cl::Buffer(
    context_,
    CL_MEM_READ_ONLY | CL_MEM_COPY_HOST_PTR,
    sizeof(cl_float) * in1.size(),
    (void*)&in1[0]);
```

- Create a buffer from a STL vector

```
kernel_.setArg(0, buf_in1_);
kernel_.setArg(1, buf_in2_);
kernel_.setArg(2, buf_out_);
```

- Set the arguments of the kernel

Host code (4)

```
queue_.enqueueNDRangeKernel(  
    kernel_,  
    cl::NullRange,  
    cl::NDRange(vector_size),  
    cl::NullRange);
```

- Execute the kernel


```
queue_.enqueueReadBuffer(  
    buf_out_,  
    CL_TRUE,  
    0,  
    vector_size * sizeof(float),  
    &out[0]);
```

- Get the result

Floyd Warshall

For each values of k (in $0..N - 1$)

The Matrix




mat (whole)

$in_x = y[k]$

$in_y = x[k]$

The Kernel Stack



Conclusion

- Third part (conclusion)
 - Optimisation tips
 - Questions?